

b

Decode latency

Matched CUDA Graph · batch 1

1.22× vs FP16

